

Roadmap for Agentic AI

1. Python & Core Programming

| Skill | Why It's Needed | Focus / Hands-On |
|---|--|---|
| Advanced Python (OOP, decorators, context managers) | Core language for AI agents; modular & maintainable code | Build utility libraries, reusable modules |
| Async Programming (<code>asyncio</code>) | Handle multiple tasks or API calls concurrently | Implement async data fetchers, concurrent agent calls |
| Logging & Exception Handling | Debugging and monitoring production agents | Add structured logging to pipelines |
| Type Hints & Static Analysis | Reduce runtime errors and improve maintainability | Use <code>mypy</code> and <code>pylint</code> |
| WebSockets (FastAPI / Socket.IO) | Real-time agent communication (streaming LLM output) | Build streaming chat demo with LLM tokens |
| File I/O & Config Management | Agents need to read/write config, models, documents | Use <code>.yaml</code> or <code>.env</code> for configs |
| Git/GitHub | Version control, collaboration | Branching, PR workflow, commit discipline |

2. Math, Stats & Machine Learning

| Skill | Why It's Needed | Focus / Hands-On |
|-------|-----------------|------------------|
|-------|-----------------|------------------|

| | | |
|---|--|--|
| Linear Algebra & Calculus | Understanding embeddings, gradients, attention mechanisms | Implement simple linear algebra operations in Python |
| Probability & Statistics | Essential for model evaluation, uncertainty, RAG retrieval scoring | Build notebooks for probability experiments & hypothesis tests |
| ML Fundamentals (Supervised/Unsupervised) | Core for embeddings, classifiers, similarity searches | Train small sklearn pipelines, evaluate metrics |
| Feature Engineering & Evaluation | High-quality inputs = better agent output | Build feature pipelines & test embeddings |

3. Deep Learning & LLMs

| Skill | Why It's Needed | Focus / Hands-On |
|--|--|--|
| Neural Networks (CNN, RNN, Transformers) | Underlying technology for embeddings & LLMs | Implement small Transformer models in PyTorch |
| Embeddings | Convert text/data into vector space for search | Use <code>sentence-transformers</code> , visualize vectors |
| Transformers & Attention | Understand how LLMs generate context-aware outputs | Inspect attention matrices, fine-tune small models |
| Prompt Engineering (Zero/Few-shot, CoT, ReAct) | Agents act and reason based on prompts | Experiment with prompts for reasoning tasks |

| | | |
|--------------------------------------|---|---|
| RAG (Retrieval-Augmented Generation) | Ground agent responses in knowledge, reduce hallucination | Build document Q&A bot using FAISS/Pinecone |
|--------------------------------------|---|---|

4. Agentic AI Systems

| Skill | Why It's Needed | Focus / Hands-On |
|--|--|--|
| Multi-Agent Systems (Planner, Worker, Committee) | Complex tasks decomposed across agents | Implement planner-worker agent to solve multi-step problems |
| Memory Management (Short-term, Long-term) | Maintain context for multi-turn interactions | Store conversation in vector DB and recall relevant info |
| LangChain / LlamaIndex | Frameworks to orchestrate agents | Build chain of tools → memory → LLMs |
| Agent Orchestration & Tool Integration | Agents need to call APIs, DBs, tools | Integrate web search, calculator, or knowledge DB into agent |

5. Backend & Production Engineering

| Skill | Why It's Needed | Focus / Hands-On |
|---------------------------|--|--|
| REST APIs (FastAPI/Flask) | Agents exposed as services | Create endpoints to query agents |
| WebSockets | Real-time streaming outputs | Stream LLM tokens to frontend for chat |
| Docker & Docker Compose | Containerized agents for reproducible environments | Containerize agent pipelines |

| | | |
|-------------------------------------|--|---|
| CI/CD (GitHub Actions/Jenkins) | Automated builds, tests, deployment | Build pipeline for ML + agents |
| Caching & Performance Optimization | Reduce LLM API cost, improve response time | Use Redis/memory caching for frequent queries |
| Logging, Monitoring & Observability | Track agent performance and failures | Integrate OpenTelemetry + Grafana dashboards |

6. Cloud & Deployment

- AWS Core Services: EC2, S3, IAM, Lambda, API Gateway
- Serverless: Deploy lightweight agent functions
- Bedrock (if available) for managed LLMs
- Infrastructure as Code (Terraform basics)
- Scaling & Load Balancing

7. Security & Reliability

- Prompt injection & guardrails
- Input validation & sanitization
- Secrets management (AWS Secrets Manager / Vault)
- Testing: unit, integration, load testing

8. Delivery & Portfolio

- Documentation (README, runbooks)
- Architecture diagrams (Draw.io / Mermaid)
- Full-stack agentic AI platform (RAG + Multi-agent + Frontend + Deployment)

Tip: Tick off each skill as you study and implement. This roadmap ensures you're fully confident and industry-ready for building agentic AI platforms.

Author - Reetesh V (Ai Engineer)

[Portfolio](#)